Fig 4 Does any implemented rule beat selecting the validation−best candidate?
negative = worse on the held−out test set; bars are 95% CI (unit = cohort)
fixed Ridge
oracle (test−set max)
descriptive reference (not implementable / post−hoc)
no selection
tree candidates
penalised
worst−half
majority vote
runner−up
argmax_val
argmax (no ref)
linear candidates
top2 ensemble
top3 ensemble
all ensemble
top10 ensemble
top5 ensemble (P2)
implemented rule
descriptive reference
pre−specified P2
−0.04
−0.02
0.00
0.02
0.04
paired difference in held−out test C−index vs argmax_val